

Course 2: Supervised Learning



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

Last session

- 1 What is not AI ?
- 2 AI definition
- 3 Applications
- 4 Open issues

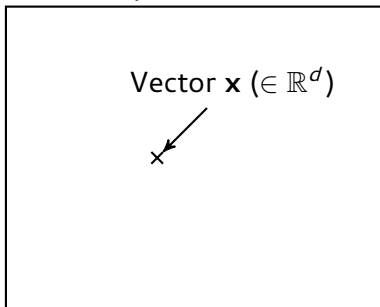
Today's session

- Learning from labeled examples
- Challenges of supervised learning

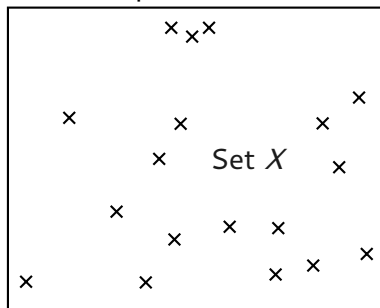
Vector space (\mathbb{R}^d)



Vector space (\mathbb{R}^d)



Vector space (\mathbb{R}^d)



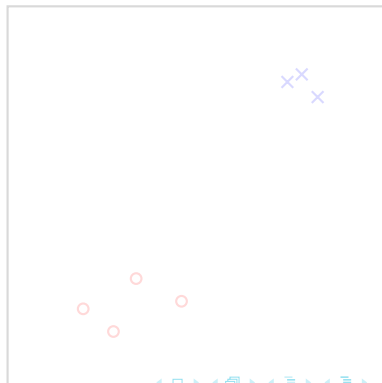
Supervised learning

Definition

Supervised learning methods use **labels** y associated to examples $x \in X$ to learn a function f such as $y = f(X)$, with the aim of **generalizing** (\neq memorizing) to unlabeled examples.

Examples

- Regression (y is scalar)
- Classification (y is categorical)
- Tons of applications:
 - Pattern recognition,
 - Prediction...



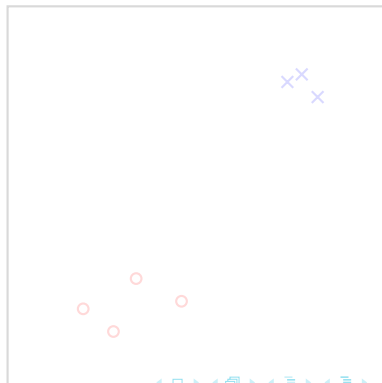
Supervised learning

Definition

Supervised learning methods use **labels** y associated to examples $x \in X$ to learn a function f such as $y = f(X)$, with the aim of **generalizing** (\neq memorizing) to unlabeled examples.

Examples

- Regression (y is scalar)
- Classification (y is categorical)
- Tons of applications:
 - Pattern recognition,
 - Prediction...



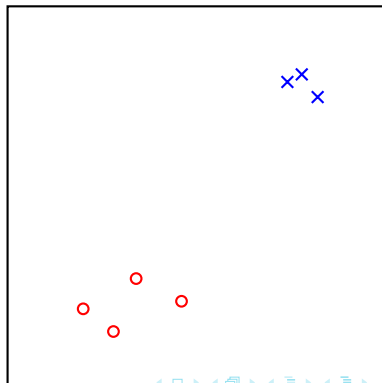
Supervised learning

Definition

Supervised learning methods use **labels** y associated to examples $x \in X$ to learn a function f such as $y = f(X)$, with the aim of **generalizing** (\neq memorizing) to unlabeled examples.

Examples

- Regression (y is scalar)
- Classification (y is categorical)
- Tons of applications:
 - Pattern recognition,
 - Prediction...



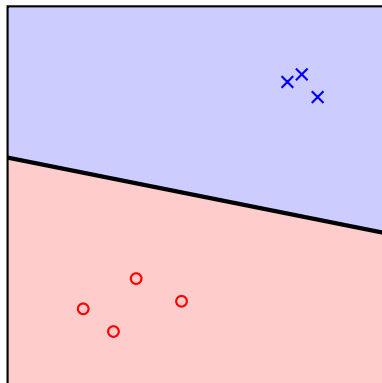
Supervised learning

Definition

Supervised learning methods use **labels** y associated to examples $x \in X$ to learn a function f such as $y = f(X)$, with the aim of **generalizing** (\neq memorizing) to unlabeled examples.

Examples

- Regression (y is scalar)
- Classification (y is categorical)
- Tons of applications:
 - Pattern recognition,
 - Prediction...



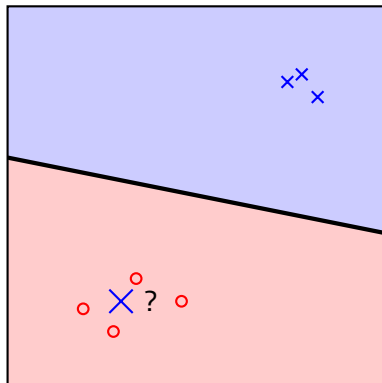
Supervised learning

Definition

Supervised learning methods use **labels** y associated to examples $x \in X$ to learn a function f such as $y = f(X)$, with the aim of **generalizing** (\neq memorizing) to unlabeled examples.

Examples

- Regression (y is scalar)
- Classification (y is categorical)
- Tons of applications:
 - Pattern recognition,
 - Prediction...



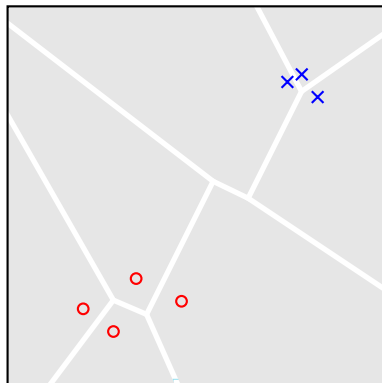
Supervised learning

Definition

Supervised learning methods use **labels** y associated to examples $x \in X$ to learn a function f such as $y = f(X)$, with the aim of **generalizing** (\neq memorizing) to unlabeled examples.

Examples

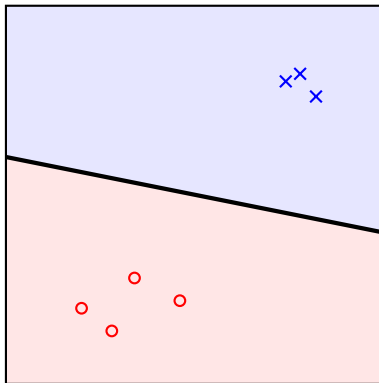
- Regression (y is scalar)
- Classification (y is categorical)
- Tons of applications:
 - Pattern recognition,
 - Prediction...



Challenges of supervised learning (1/5)

An ill-defined problem

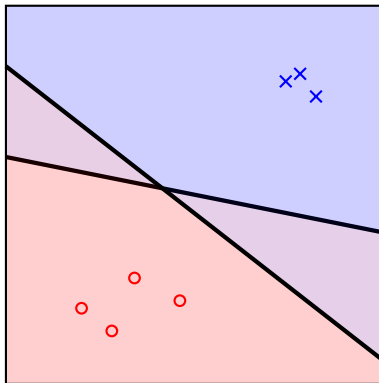
- An infinity of potential solutions, one must be the “best one” but is unreachable,
- \Rightarrow requires a priori, constraints.



Challenges of supervised learning (1/5)

An ill-defined problem

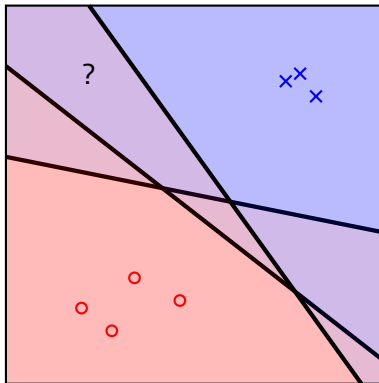
- An infinity of potential solutions, one must be the “best one” but is unreachable,
- \Rightarrow requires a priori, constraints.



Challenges of supervised learning (1/5)

An ill-defined problem

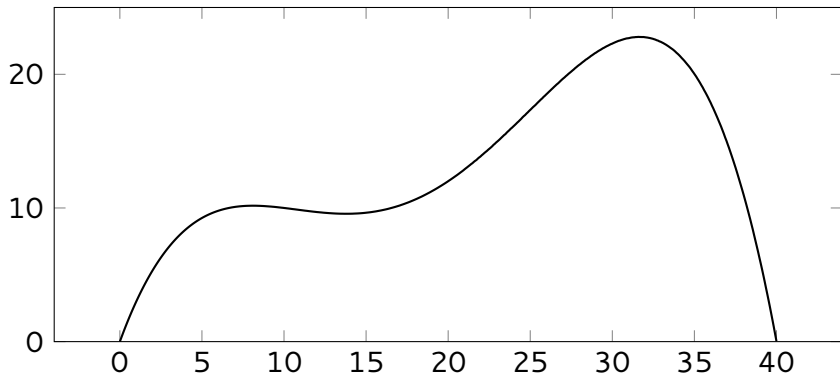
- An infinity of potential solutions, one must be the “best one” but is unreachable,
- \Rightarrow requires a priori, constraints.



Challenges of supervised learning (2/5)

Bias/variance trade-off

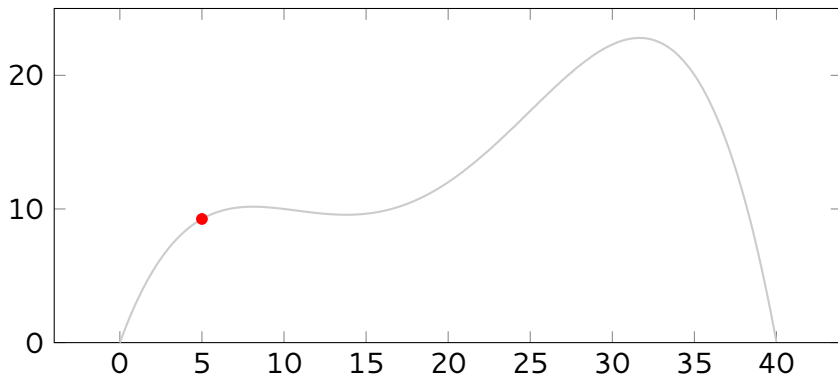
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

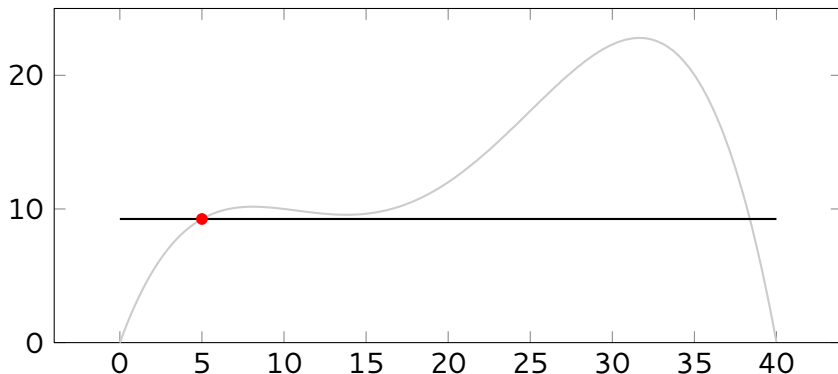
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

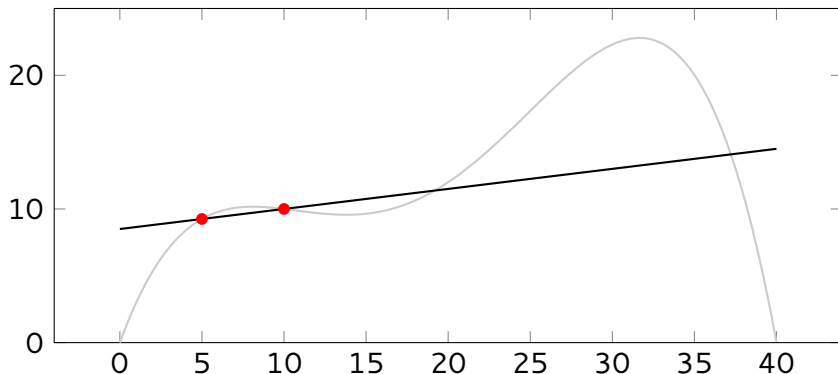
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

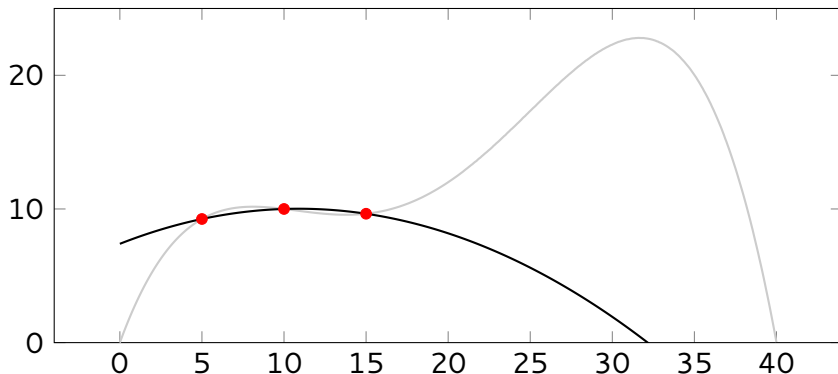
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

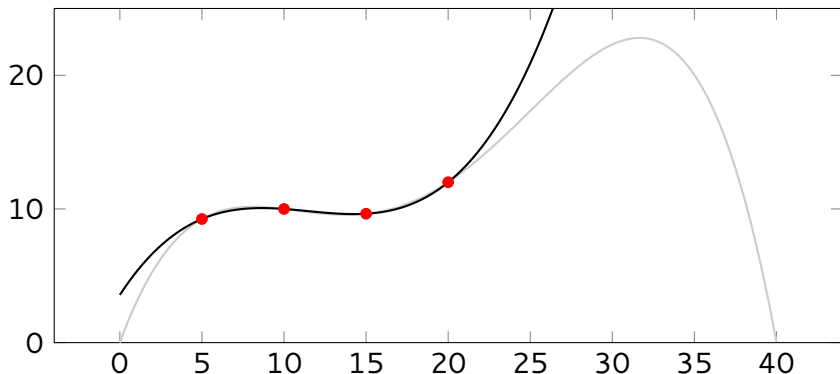
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

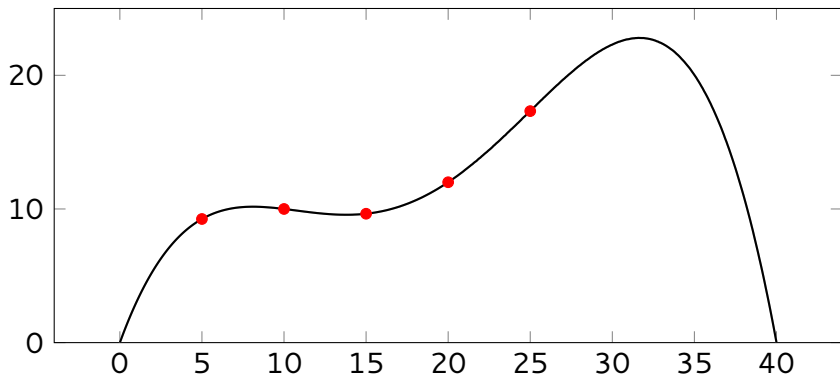
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

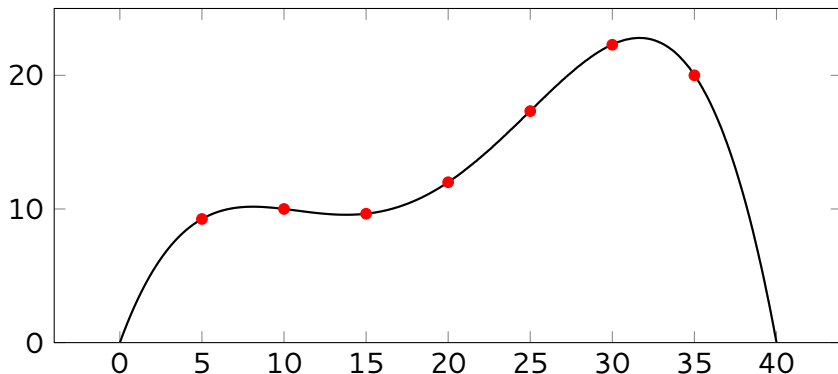
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

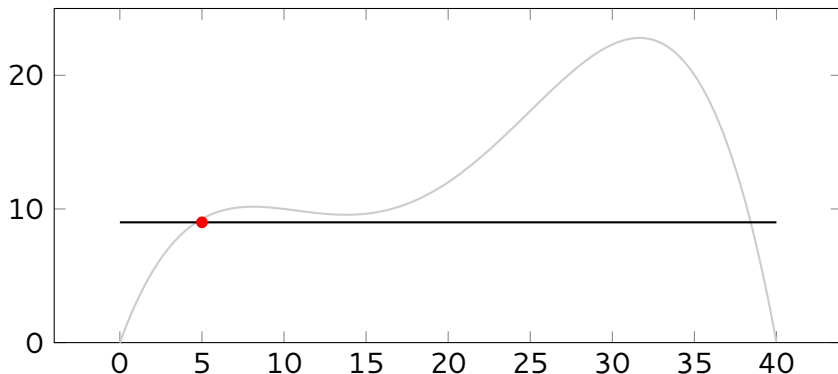
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

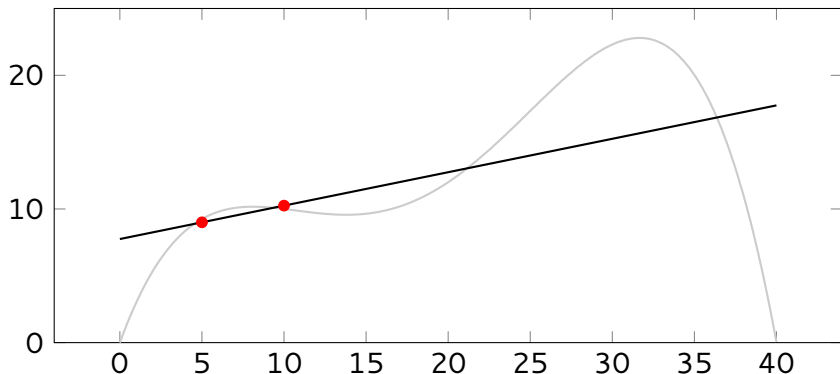
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

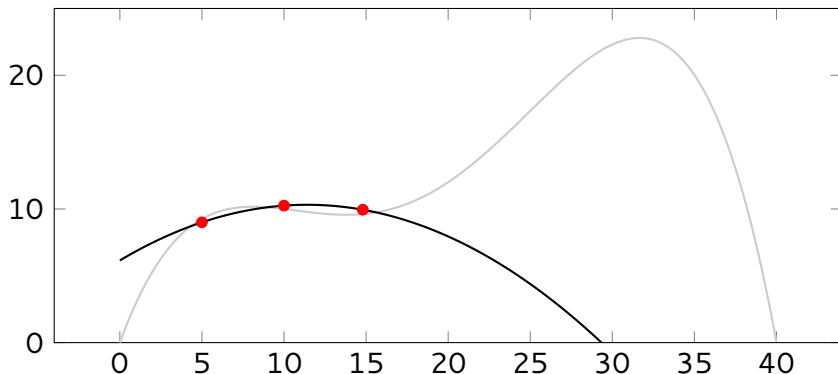
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

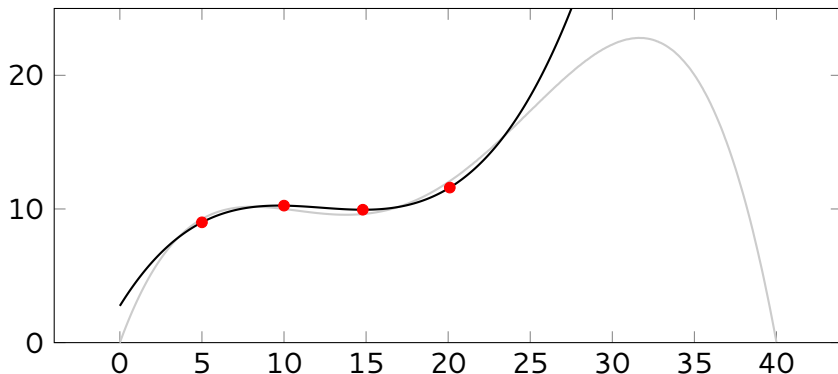
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

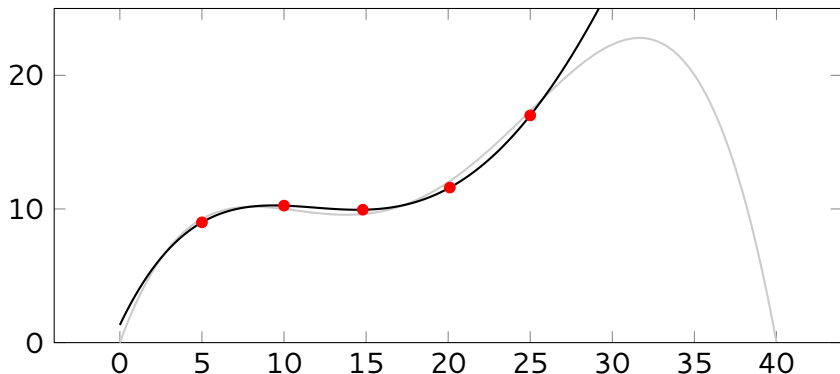
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

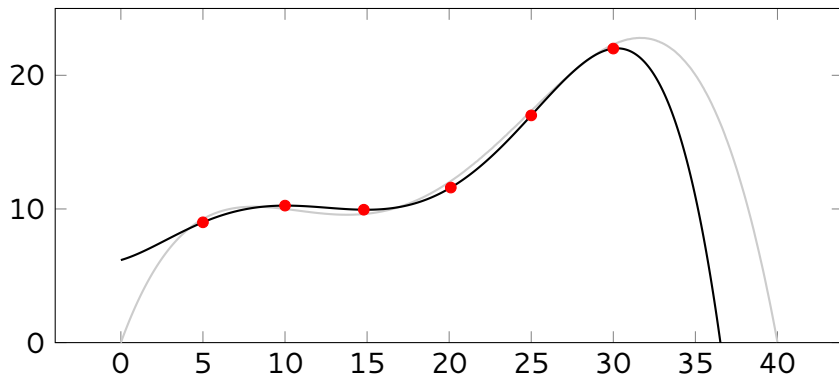
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

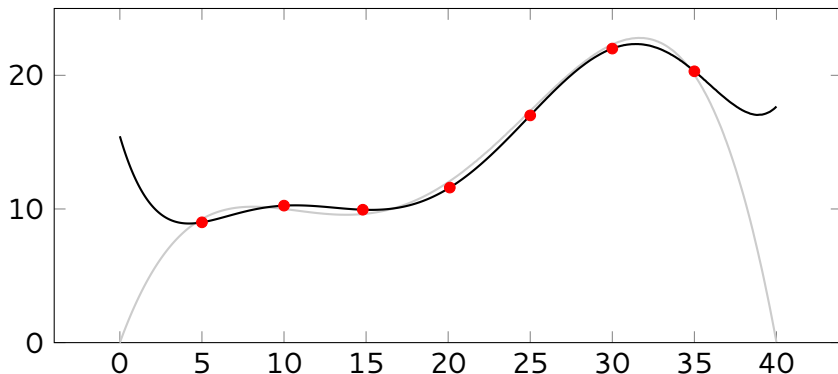
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

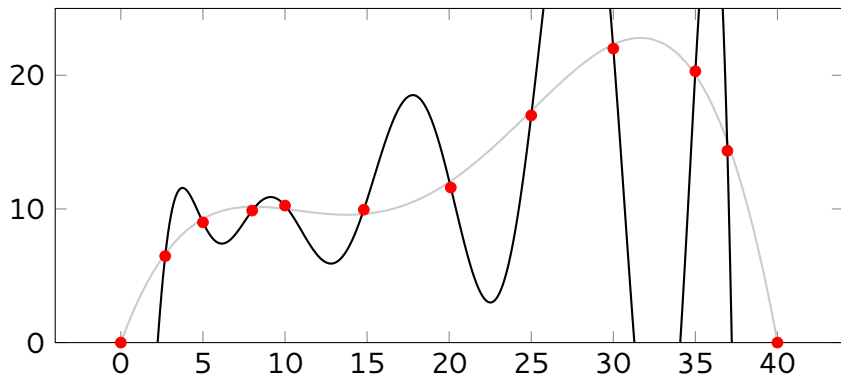
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

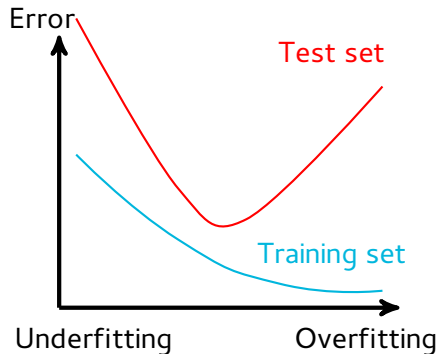
- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



Challenges of supervised learning (2/5)

Bias/variance trade-off

- A simple solution that almost matches is better than a complex one that fully matches,
- Mimicking is not learning: overfitting problem.



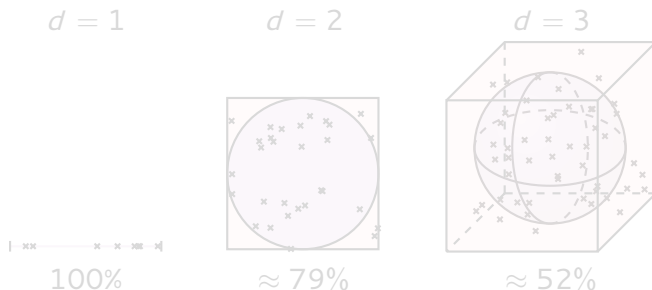
Crossvalidation

- To quantify overfitting, split training dataset in two parts:
 - 1 A first part is used to train,
 - 2 A second part is used to validate,

Challenges of supervised learning (3/5)

Curse of dimensionality

- Geometry is not intuitive in high dimension,
- Efficient methods in 2D are not necessarily still valid.



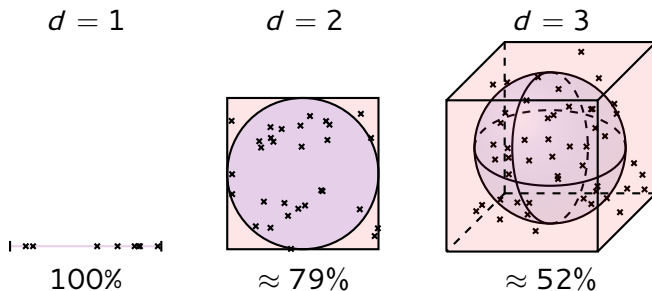
$$V_d^s = \frac{\pi^{d/2} R^d}{\Gamma(d/2 + 1)} \text{ versus } V_d^c = (2R)^d$$

see <http://www.maths.manchester.ac.uk/~mlotz/teaching/suprises.pdf>

Challenges of supervised learning (3/5)

Curse of dimensionality

- Geometry is not intuitive in high dimension,
- Efficient methods in 2D are not necessarily still valid.

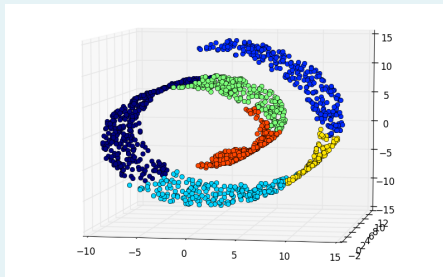


$$V_d^S = \frac{\pi^{d/2} R^d}{\Gamma(d/2 + 1)} \text{ versus } V_d^C = (2R)^d$$

see <http://www.maths.manchester.ac.uk/~mlotz/teaching/surprises.pdf>

Challenges of supervised learning (4/5)

Riemannian manifolds

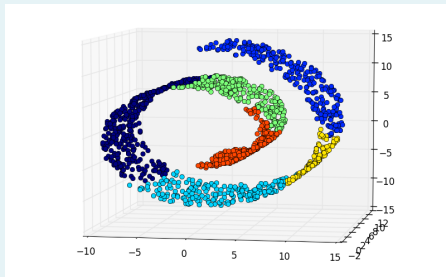


Linear separability and need for embedding

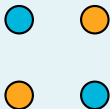


Challenges of supervised learning (4/5)

Riemannian manifolds



Linear separability and need for embedding



Challenges of supervised learning (5/5)

Computation time

Example on ImageNet, simply going through all images:

- $n = 10.000.000$, $d \approx 1.000.000$,
- $\approx 10^{13}$ elementary operations,
- $\approx 2\text{h}45$ on a modern processor.

Scalability

- Finding the best solution to a problem would be feasible with unlimited computation time,
- But searching through the space of possible functions is often untractable,
- Solutions must be computationally reasonable, which is the true challenge today.

Challenges of supervised learning (5/5)

Computation time

Example on ImageNet, simply going through all images:

- $n = 10.000.000$, $d \approx 1.000.000$,
- $\approx 10^{13}$ elementary operations,
- $\approx 2\text{h}45$ on a modern processor.

Scalability

- Finding the best solution to a problem would be feasible with unlimited computation time,
- But searching through the space of possible functions is often untractable,
- Solutions must be computationally reasonable, which is the true challenge today.

Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.

Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.

x x

Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.

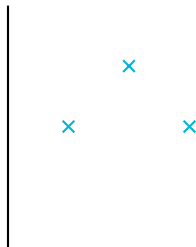


Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.

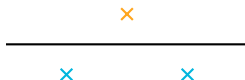


Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



Vapnik Chervonenki (VC) dimension

Definition

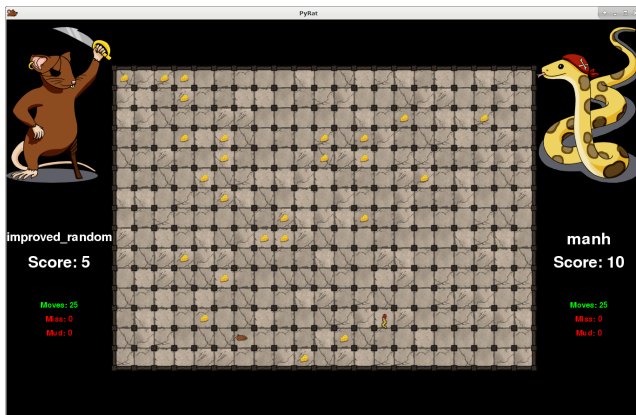
- Let us fix d ,
- The VC dimension is a measure of the genericity of a method,
- It is the maximum cardinality of a set of vectors that the method is able to shatter in any possible way.

Consider for example lines to shatter set of points with $d = 2$.



VC is 3.

Non-symmetric PyRat without walls / mud



Both players follow a deterministic greedy algorithm.
Supervised learning - predict the outcome of a game from the start configuration.
Expected accuracy of a random classifier ?

Lab Session 2 and assignments for Session 3

TP Supervised Learning (TP1)

- Basics of machine learning using sklearn (including new definitions / concepts)
- Generating PyRat Datasets
- Tests on PyRat datasets using a naive approach

Project 1 (P1)

You will be assigned a supervised learning method. You have to prepare a Jupyter Notebook on this method, including:

- A brief description of the theory behind the method,
- Basic tests on simulated data to show the influence of parameters and hyperparameters
- Advanced tests and analysis on your own PyRat Datasets (including changing the dataset)

During Session 3 you will have 7 minutes to present your notebook.